

Exercises with special data structures

1. Given an array containing n integers, show how to find the largest k elements in the array in $O(n + k \log n)$ time. Given k arrays which are individually sorted in non-decreasing order, show how to merge them into a single sorted array in $O(n \log k)$ time, where n is the total number of elements in all the arrays.

2. Given n tasks such that the i th task has a release time r_i , execution time t_i and deadline d_i . The task can start execution at any time $\geq r_i$, and must finish at or before deadline d_i . There is only one machine available and it can execute at most one task at a time. Other tasks can be kept pending. It is possible to switch from one task to another without any loss of time. A task must be executed for t_i time, but it can be interrupted any number of times. Describe an $O(n \log n)$ time algorithm to determine whether it is possible to complete all tasks before their deadlines.

3. The suffix tree of a string S of length n is a compressed trie that contains all the $n + 1$ suffixes of the string. Usually a distinct $\$$ symbol is added at the end of the string, so that no suffix is a prefix of another. Each leaf node of the tree is labelled by the starting index of the corresponding suffix. A compressed trie means that nodes that have only one child are eliminated and edges are labelled by substrings rather than single characters. The label is represented by a pair of integers giving the starting position and length of the substring. This reduces the space requirement to $O(n)$. A slightly easier structure to construct is the suffix array, which is a permutation of indices $0, \dots, n$ in the sorted order of the suffixes. Given the suffix tree, show how the suffix array can be constructed in $O(n)$ time. Constructing the suffix tree in $O(n)$ time is non-trivial. Can you construct the suffix array directly in $O(n)$ time? Another useful array is the lcp array, that gives the longest common prefix between the i th and $(i + 1)$ th suffix in sorted order. Show how this can be constructed in $O(n)$ time, from the suffix tree, and also directly.

Given the suffix tree, show how you can test whether a given pattern string p is a substring of S in time $O(p)$. Show how to compute the number of different substrings of S from the suffix tree in $O(n)$ time. Show how to do this using the suffix and lcp arrays, with an extra $O(\log n)$ time in the case of pattern matching.

4. A hash function can be used to check for errors in a file that may occur during transmission. Assuming the file is a bit string s , a hash value $h(s)$ is transmitted along with s , and recomputed for the received file. If the values mismatch, it indicates an error. It is assumed there is no error in the hash value received. Describe a simple hash function h such that if the received file differs in exactly one bit, the hash value is guaranteed to be different. Suppose a hash function is required that can detect errors in at most 2 bits. Describe such a hash function. What is the minimum number of bits required in the hash value? Can you get a function that achieves the minimum? One hash function that is actually used is called the md5 check-sum.